



INGENIOUS 7.0

HACKATHON

Team Details

- Team name: **Smooth_Operator**
- Team leader name: **Het Modi**
- Team leader email-id: **m.hetmodi@gmail.com**
- Problem Statement: **Cyber-Resilient Infrastructure for Critical Sectors**
- Team Members: **Chinmay Patel, Aryan Vadhadiya, Jay Shah, Het Patel**

Problem Overview

Critical digital infrastructures in healthcare, agriculture, and urban systems operate at large scale and are expected to be always available and secure. However, these **systems are increasingly exposed to sophisticated cyber threats**, many of which operate over encrypted traffic and evade traditional security mechanisms.

Our Idea

- An intelligent, behaviour-based security system that **detects zero-day and advanced threats hidden in encrypted traffic without decryption**.
- **Machine learning-driven** anomaly detection enables automated, **rapid threat containment** and delivers an adaptive, **future-ready cyber-resilience layer** for critical infrastructure.

Our Progress

- Implemented kubernetes with horizontal pod scaling.
- Add email on alert/attack functionality.
- Added rate limiting (customisable value) in the node software.
- Added ip blocking in by node.
- Developed a unified setup utility that provisions the complete Kubernetes environment.
- Federated learning on telemetry data.

Proof of Concept

| Problem → Solution Mapping | | |
|--------------------------------------|--|--------|
| Problem Requirement | Your Solution | Status |
| Real-time monitoring | Ingest Service + WebSocket | ✓ |
| Anomaly detection algorithms | ML Models (LSTM, RF) + Rule Engine | ✓ |
| Alerting system with severity levels | Alert Manager (low/medium/high/critical) | ✓ |
| Automated response mechanisms | Response Engine + Playbooks | ✓ |
| Dashboard showing security status | React Frontend | ✓ |
| Simulated attack scenarios | Attack Simulator Scripts | ✓ |
| Healthcare/Agriculture/Urban domains | Domain-specific ML models | ✓ |
| Scalability | Kubernetes + HPA | ✓ |



Approach

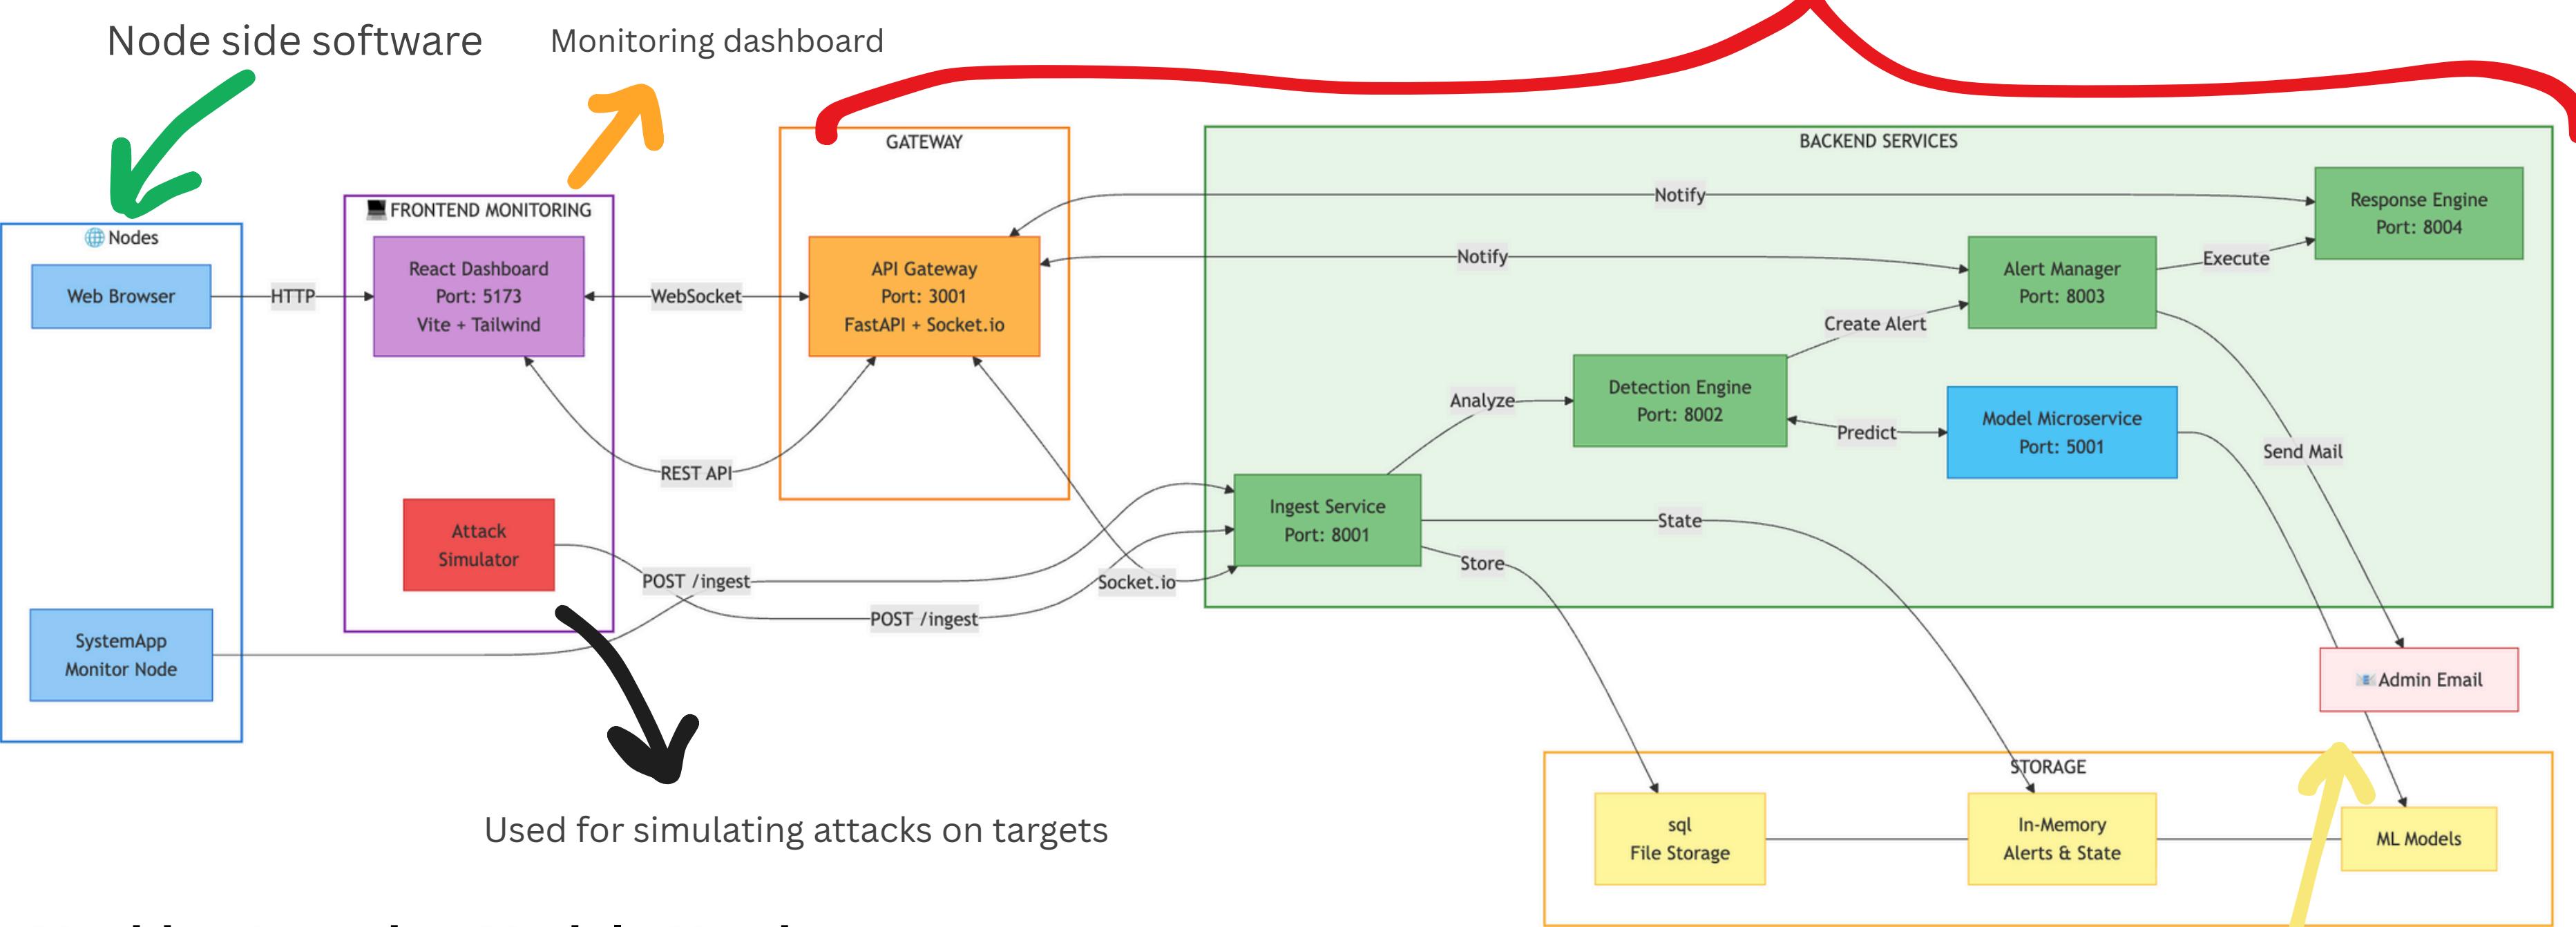
- Traffic Capture & Flow Construction
- **Network traffic is captured at the edge**, extracting timing, volume, and protocol-level metadata.
- Encryption-Aware Routing
- **Rule-Based Analysis for Unencrypted Traffic**
- Unencrypted traffic is inspected using predefined signatures and security rules for fast, deterministic detection.
- Behaviour-Based Analysis for Encrypted Traffic
- **Encrypted traffic bypasses payload inspection and is analysed using flow-level behavioural features.**
- Domain-Aware Model Selection
- The system dynamically selects the appropriate detection model based on the operational context:
 - General Network → **Neural Network Classifier**
 - Agriculture → **LSTM Autoencoder**
 - Healthcare → **LSTM Classifier**
 - Urban → **LSTM Forecaster**
- Threat Scoring & Decision Engine
- Automated & Manual Response Execution
- **Federated Learning Feedback Loop**
- **Local learnings are aggregated centrally and redistributed,**

Architecture Diagram

Each service has its own **docker** image which is then scaled on **kubernetes**

Node side software

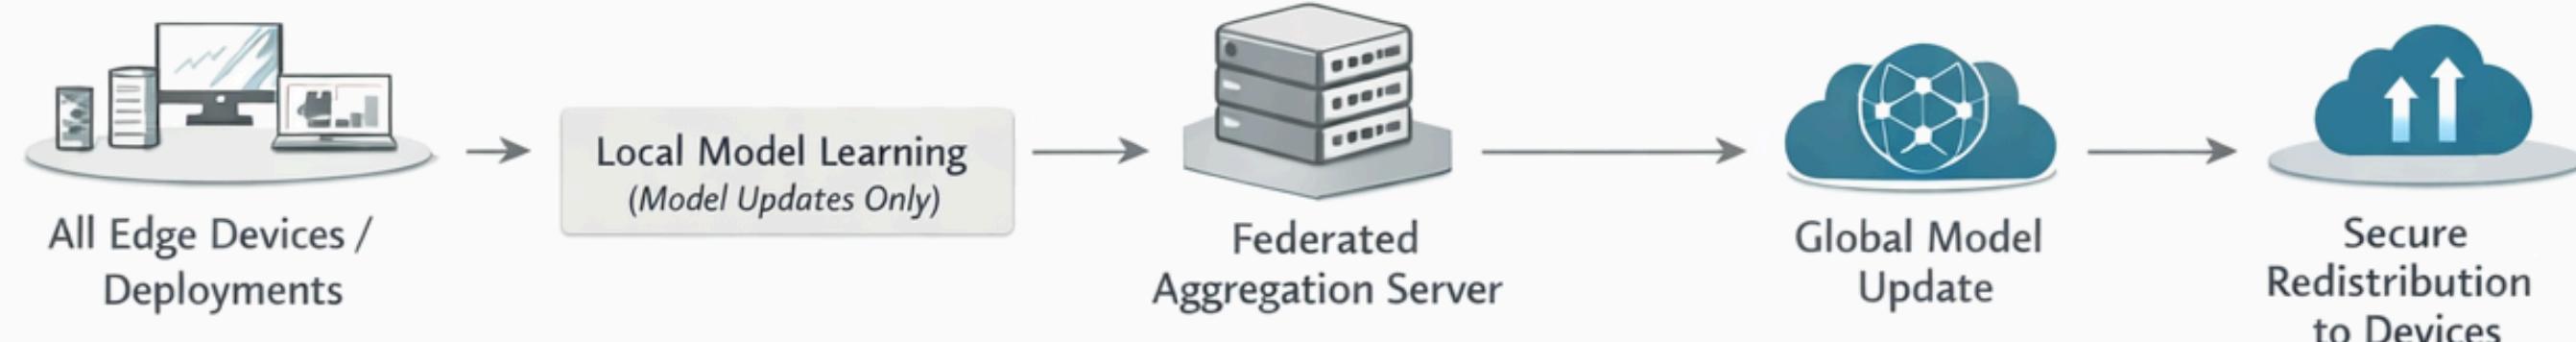
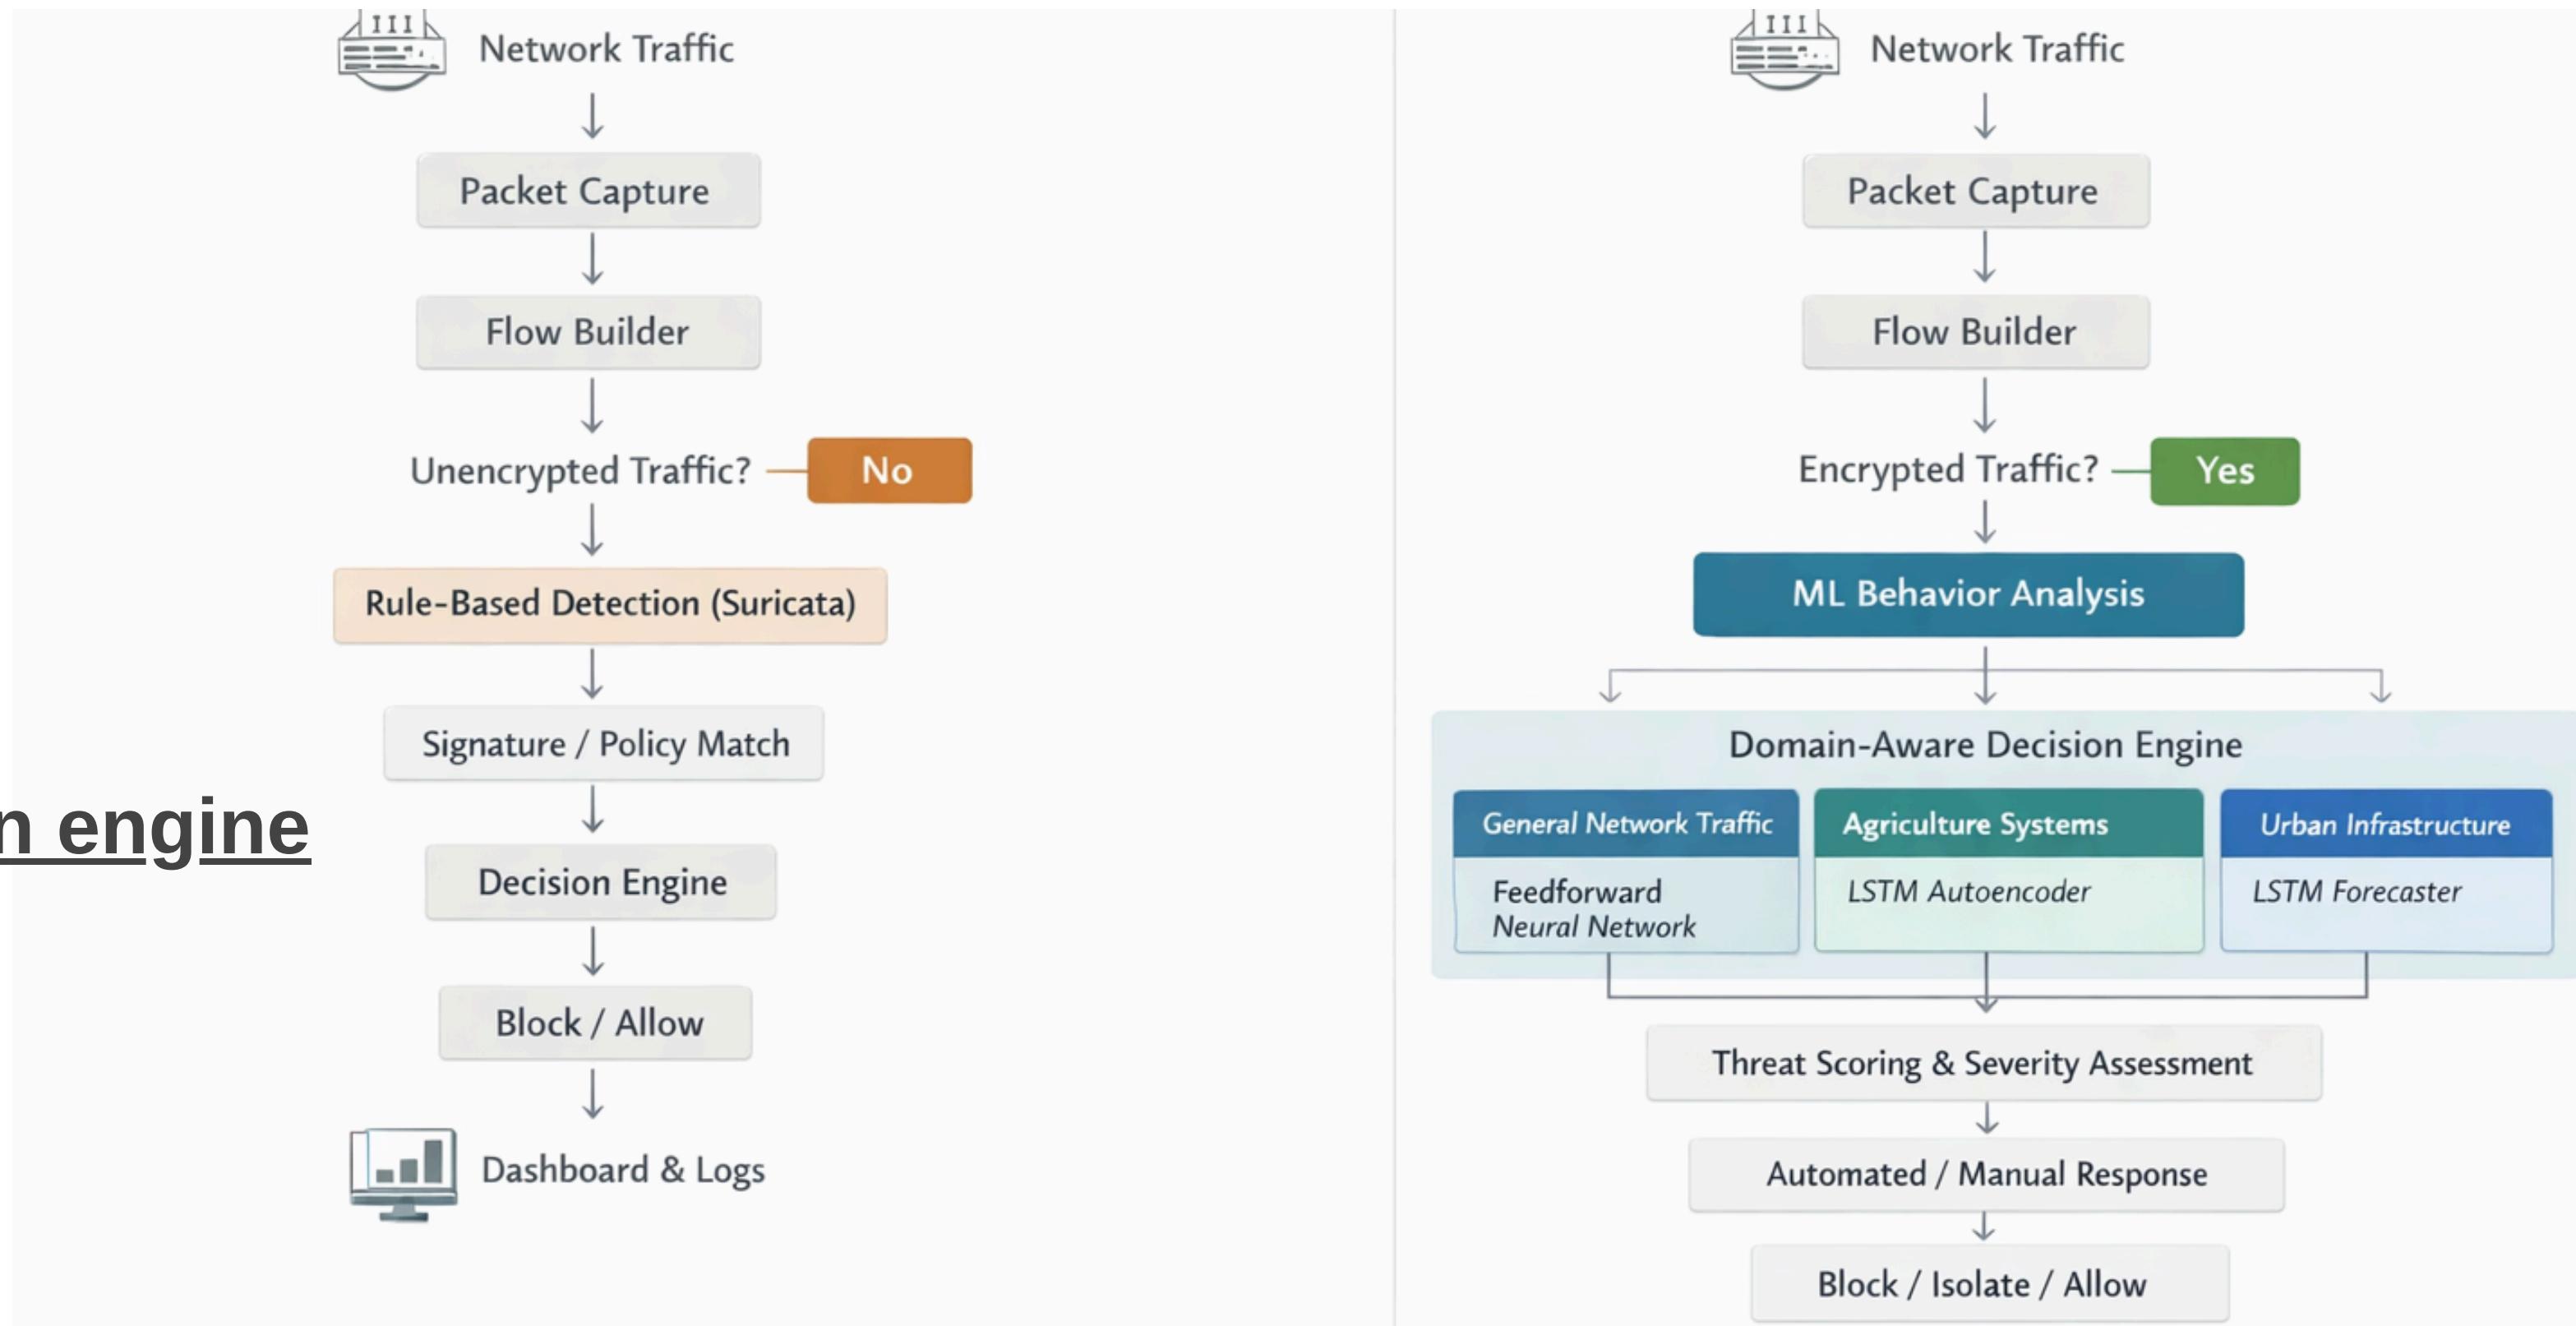
Monitoring dashboard



Machine Learning Models Used

1. **Feedforward Neural Network (PyTorch)**: **Generic network attack detection** (DDoS, Botnet, Brute Force)
2. **Isolation Forest (Unsupervised ML)**: Network flow based anomaly detection on encrypted traffic
3. **LSTM Autoencoder (Unsupervised)**: Agriculture sensor spoofing and physics-based anomaly detection
4. **LSTM Classifier (Supervised)**: Healthcare IoMT DDoS and traffic surge detection
5. **LSTM Forecaster (Regression)** : Urban traffic flow prediction and grid manipulation detection

Detection engine



Scalability: Kubernetes

- The platform uses Kubernetes **Horizontal Pod Autoscaler** to dynamically scale services based on real-time resource utilization.
- **Self-Healing Infrastructure:** Failed or compromised pods are automatically recreated.
- **Multi-Replica Critical Services:** Detection, alerting, and dashboard services run with replicas.

```
(base) jay@Jays-MacBook-Air Smooth_Operator % ./k8s/kuber_start.sh
Starting Threat Ops Kubernetes Deployment...
✓ Kubernetes cluster detected
✓ Docker images found
部署 to Kubernetes...
deployment.apps/ingest-service configured
service/ingest-service unchanged
deployment.apps/api-gateway configured
service/api-gateway unchanged
deployment.apps/detection-engine configured
service/detection-engine unchanged
deployment.apps/alert-manager configured
service/alert-manager unchanged
deployment.apps/response-engine configured
service/response-engine unchanged
deployment.apps/frontend configured
service/frontend unchanged
deployment.apps/model-microservice configured
service/model-microservice unchanged
deployment.apps/systemapp configured
service/systemapp unchanged
ingress.networking.k8s.io/threat-ops-ingress unchanged
```

| Name | Images | Labels | Node | Status | Restarts | CPU Usage (cores) | Memory Usage (bytes) | Created | ⋮ |
|--|--|---|----------------|---------|----------|-------------------|----------------------|-------------|---|
| api-gateway-85466964b5-5x9jd | api-gateway:latest | app: api-gateway pod-template-hash: 85466964b5 | docker-desktop | Running | 0 | 3.00m | 42.52Mi | an hour ago | ⋮ |
| alert-manager-57bb875bd9-kvzb6 | alert-manager:latest | app: alert-manager pod-template-hash: 57bb875bd9 | docker-desktop | Running | 0 | 2.00m | 40.52Mi | an hour ago | ⋮ |
| ingest-service-64cc9599cc-62bct | ingest-service:latest | app: ingest-service pod-template-hash: 64cc9599cc | docker-desktop | Running | 0 | 4.00m | 44.50Mi | an hour ago | ⋮ |
| ingest-service-64cc9599cc-tkwr6 | ingest-service:latest | app: ingest-service pod-template-hash: 64cc9599cc | docker-desktop | Running | 0 | 23.00m | 46.19Mi | an hour ago | ⋮ |
| kubernetes-dashboard-79cbf9fb6-94lrg | kubernetesui/dashboard:v2.7.0 | k8s-app: kubernetes-dashboard pod-template-hash: 79cbf9fb6 | docker-desktop | Running | 0 | 1.00m | 30.08Mi | 3 hours ago | ⋮ |
| dashboard-metrics-scraper-5bd45c9dd6-z6wwj | kubernetesui/metrics-scraper:v1.0.8 | k8s-app: dashboard-metrics-scraper pod-template-hash: 5bd45c9dd6 | docker-desktop | Running | 0 | 1.00m | 14.90Mi | 3 hours ago | ⋮ |
| metrics-server-56fb9549f4-54rms | registry.k8s.io/metrics-server/metrics-server:v0.8.0 | k8s-app: metrics-server pod-template-hash: 56fb9549f4 | docker-desktop | Running | 0 | 5.00m | 28.73Mi | 3 hours ago | ⋮ |
| ingress-nginx-controller-7bd4bbb674-7cjq2 | registry.k8s.io/ingress-nginx/controller:v1.10.1@sha256:e24f39d3eed6bcc239a56f20098878845f62baa34b9f2be2fd2c38ce9fb0f29e | app.kubernetes.io/component: controller app.kubernetes.io/instance: ingress-nginx app.kubernetes.io/name: ingress-nginx | docker-desktop | Running | 0 | 3.00m | 135.57Mi | 5 hours ago | ⋮ |
| vpnkit-controller | docker/desktop-vpnkit-controller:dc31cb22850be0cdd97c84a9cfecaf44a1afbfe | component: vpnkit-controller | docker-desktop | Running | 0 | 1.00m | 12.83Mi | 5 hours ago | ⋮ |
| storage-provisioner | docker/desktop-storage-provisioner:2.0 | component: storage-provisioner | docker-desktop | Running | 0 | 2.00m | 13.98Mi | 5 hours ago | ⋮ |

ZERO-DAY & ENCRYPTED ATTACK DETECTION

- **Live Packet Capture** : Uses **unsupervised machine learning** to learn effectively detects previously unseen attacks that bypass signature-based security systems

LAYERED AI DEFENSE ARCHITECTURE



- Ensures early-stage filtering of broad attacks while enabling **deep contextual analysis** for healthcare, agriculture, and urban systems.

○

FEDERATED LEARNING-ENABLED COLLECTIVE DEFENSE

- Only model updates and threat insights are shared and **aggregated centrally** to improve global detection accuracy.

KUBERNETES-BASED SCALABLE DEPLOYMENT

- Uses containerization and Kubernetes orchestration to automatically scale edge services and central components based on traffic load, ensuring high availability, fault tolerance, and consistent performance under attack conditions.

PROJECT FEATURES

AUTOMATED RESPONSE & POLICY ENFORCEMENT (SOAR)

- Automatically blocks malicious IPs, and escalates alerts based on threat severity and confidence levels.

REAL-TIME DASHBOARD & VISUALIZATION

- Displays **live telemetry, alert timelines, device status**, and response actions through **WebSockets**.

ATTACK SIMULATION & TESTING FRAMEWORK

- Includes **attacker and victim simulation** tools to validate detection and response workflows.

AUTOMATED ALERT GENERATION & SEVERITY CLASSIFICATION

- Transforms detected security issues into human-readable alerts through emails with clear severity levels.



CODE SNIPPETS

ANN - GENERIC NETWORK ATTACKS



```

1 client_id = data.get('client_id', 'unknown')
2         client_weights = data.get('weights')
3
4 if not client_weights:
5     print(f"[ERR] No weights in payload from {client_id}")
6     return jsonify({"error": "Missing weights"}), 400
7
8 print(f"[INFO] Processing update from {client_id} (Keys: {len(client_weights)})")
9
10 client_state_dict = {k: torch.tensor(v).to(DEVICE) for k, v in client_weights.items()}
11
12 pending_updates.append(client_state_dict)
13 print(f"[FL] Received FL update from {client_id}. Progress: {len(pending_updates)}/{REQUIRED_UPDATES_FOR_AVG}")
14
15 if len(pending_updates) >= REQUIRED_UPDATES_FOR_AVG:
16     print("[FL] Aggregating Global Model (FedAvg)...")
17
18 # FEDERATED AVERAGING ALGORITHM
19 avg_weights = copy.deepcopy(GLOBAL_MODEL.state_dict())
20 for key in avg_weights.keys():
21     # Stack all updates for this key and take mean
22     # Ensure all updates have this key
23     valid_updates = [u[key] for u in pending_updates if key in u]
24     if valid_updates:
25         stacked = torch.stack(valid_updates)
26         avg_weights[key] = torch.mean(stacked, dim=0)
27
28 # Update Global Model
29 GLOBAL_MODEL.load_state_dict(avg_weights)
30 FL_ROUND += 1
31 pending_updates = [] # Reset
32 print(f"[SUCCESS] Global Model Updated! New Round: {FL_ROUND}")

```



```

1 @app.route('/api/f1/submit-update', methods=['POST'])
2 def submit_update():

```

FEDERATED LEARNING ENDPOINT



```

1 train_idx, _ = train_test_split(range(len(tensor_x)), test_size=0.2)
2         loader = DataLoader(TensorDataset(tensor_x[train_idx], tensor_y[train_idx]), batch_size=2048, shuffle=True)
3
4 model = GeneralNetworkShield(input_dim=X.shape[1]).to(DEVICE)
5 optimizer = optim.Adam(model.parameters(), lr=0.001)
6 criterion = nn.BCELoss()
7
8 model.train()
9 for epoch in range(3):
10     total_loss = 0
11     for X_batch, y_batch in loader:
12         X_batch, y_batch = X_batch.to(DEVICE), y_batch.to(DEVICE)
13         optimizer.zero_grad()
14         out = model(X_batch)
15         loss = criterion(out, y_batch)
16         loss.backward()
17         optimizer.step()
18         total_loss += loss.item()
19     print(f" Epoch {epoch+1} Loss: {total_loss/len(loader):.4f}")

```

POST http://127.0.0.1:5000/api/analyze

Body (raw)

```
{
  "sector": "healthcare",
  "payload": "SELECT * FROM patients WHERE id=1 OR 1=1",
  "sensor_data": [60, 10, 5, 5]
}
```

Body (JSON)

```
{
  "message": "Malicious Web Payload Detected (SQLi/XSS)",
  "source": "Web Gatekeeper",
  "status": "blocked",
  "threat_level": "high"
}
```

POSTMAN MODEL TESTING

HEALTH CARE ML MODEL

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Normal | 1.00 | 0.98 | 0.99 | 47974 |
| DDoS | 0.99 | 1.00 | 0.99 | 51893 |
| accuracy | | | | 0.99 |
| macro avg | 0.99 | 0.99 | 0.99 | 99867 |
| weighted avg | 0.99 | 0.99 | 0.99 | 99867 |

Accuracy: 0.9910280673295483

Snapshot of Node software

- Agent installed on protected machines - runs as a **background daemon**.
- Auto-discovers device identity and sends **real-time telemetry** to our dashboard.
- Captures attacks while monitoring **system health**.
- **Cross-platform support** - same agent works on **Windows, Mac, and Linux**
- Supports **headless** systems, using command line arguments.

The screenshot displays two tabs of the Threat_Ops Agent Console interface, both running at 127.0.0.1:5050.

Top Tab (System Diagnostics):

Device ID: JAYS-MACBOOK-AIR.LOCAL-NODE
IP Address: 10.110.5.129

| Category | Value |
|-------------------|----------|
| CPU Usage | 23.3% |
| Memory Usage | 67% |
| Disk Usage | 35.5% |
| Uptime | 02:50:17 |
| CPU Cores/Threads | 10 / 10 |
| Total Memory | 16 GB |
| Active Processes | 709 |
| Load Average | 1.46 |
| Network Sent | 203 MB |
| Network Received | 267.1 MB |

Bottom Tab (Virtual IoT Fleet):

Ingest End Point: localhost:8001
Sector: HEALTHCARE

UPDATE TARGET button

Virtual IoT Fleet (Sector: Healthcare):

| Device | Type | Status | Health |
|------------------|---------------|-------------|--------|
| infusion-pump-03 | INFUSION PUMP | COMPROMISED | 10% |
| mri-scanner-02 | MRI SCANNER | COMPROMISED | 100% |
| pacemaker-01 | PACEMAKER | COMPROMISED | 100% |

Actions: HEAL ALL DEVICES, RESET FLEET

Snapshot of Monitoring dashboard

The screenshot shows a monitoring dashboard for the Threat_Ops.ai Security Operations Center, accessible at localhost:5173. The interface is dark-themed and displays real-time system telemetry and alerting information.

Connected Nodes: 2 Active. Nodes listed: Het Workspace Node (Online) and Jays Macbook Air.Local Node (Online).

System Telemetry: Shows CPU (12.3%), Memory (66.8%), and Network (0.0 KB/s) usage over the last 30 minutes. A line chart visualizes the memory usage, which spikes significantly around 00:02 and 00:15.

Alerts: 15 critical alerts are active. Two recent events are listed as SQL Injection Attacks detected on Jays-MacBook-Air.local-node, both marked as Critical and acknowledged 15s ago.

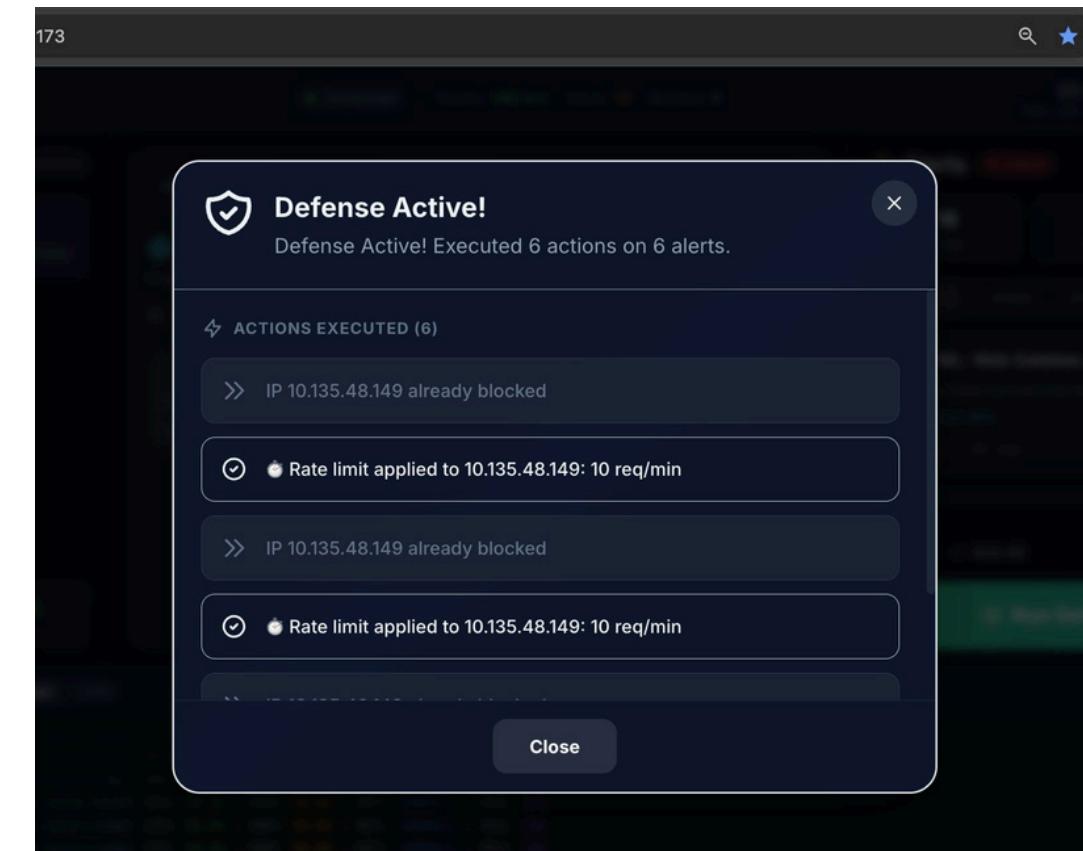
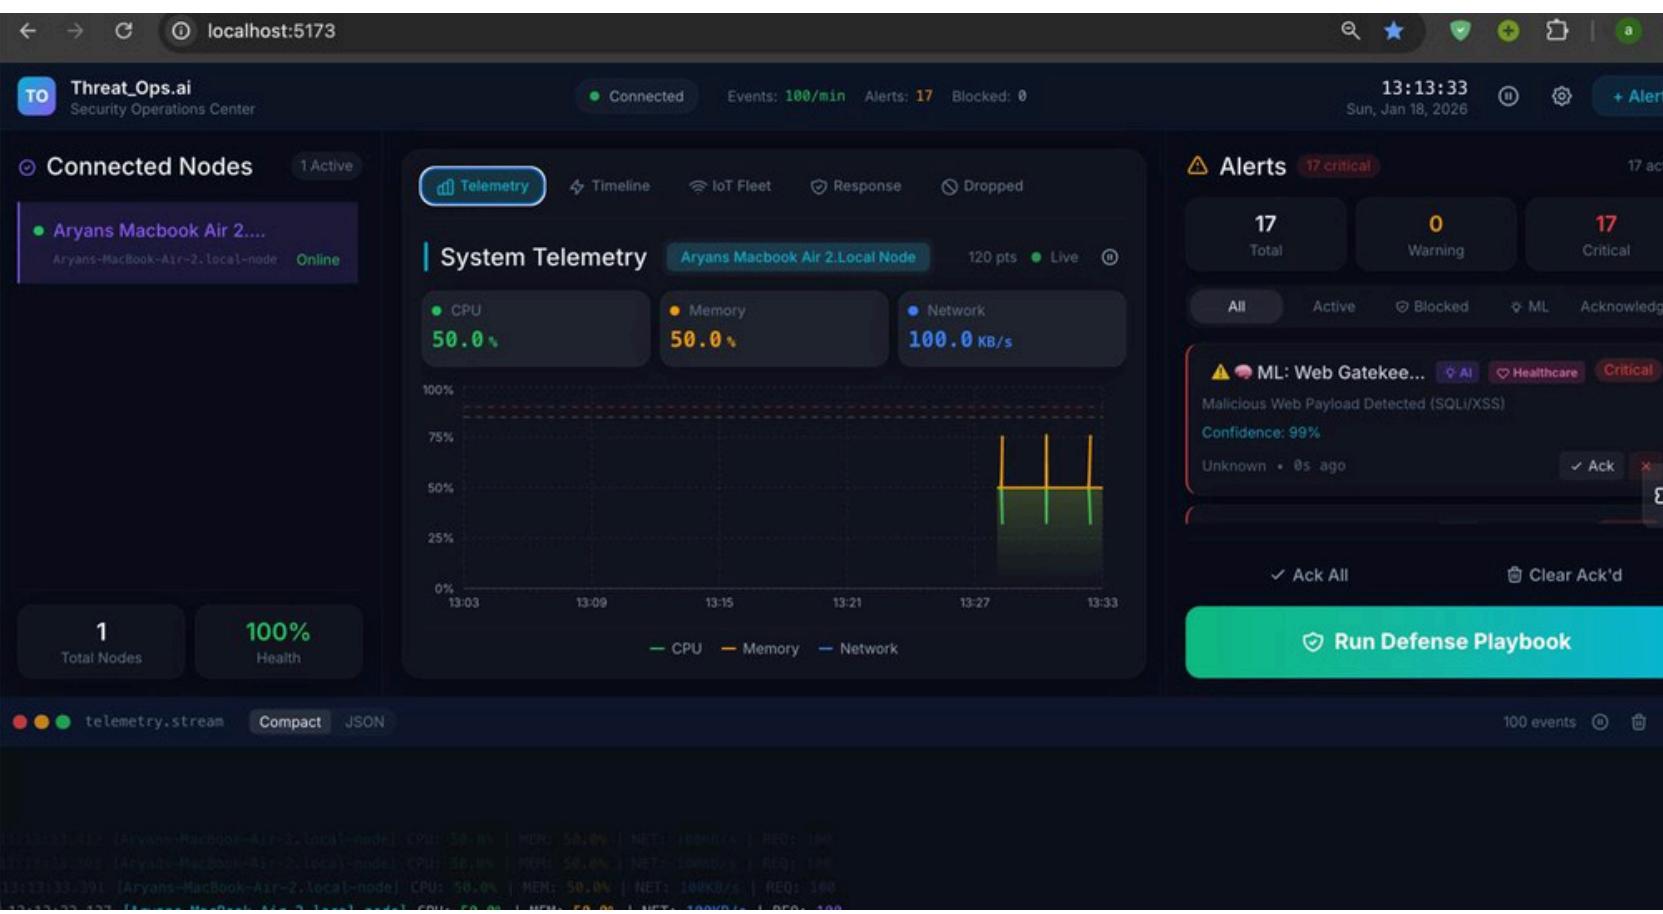
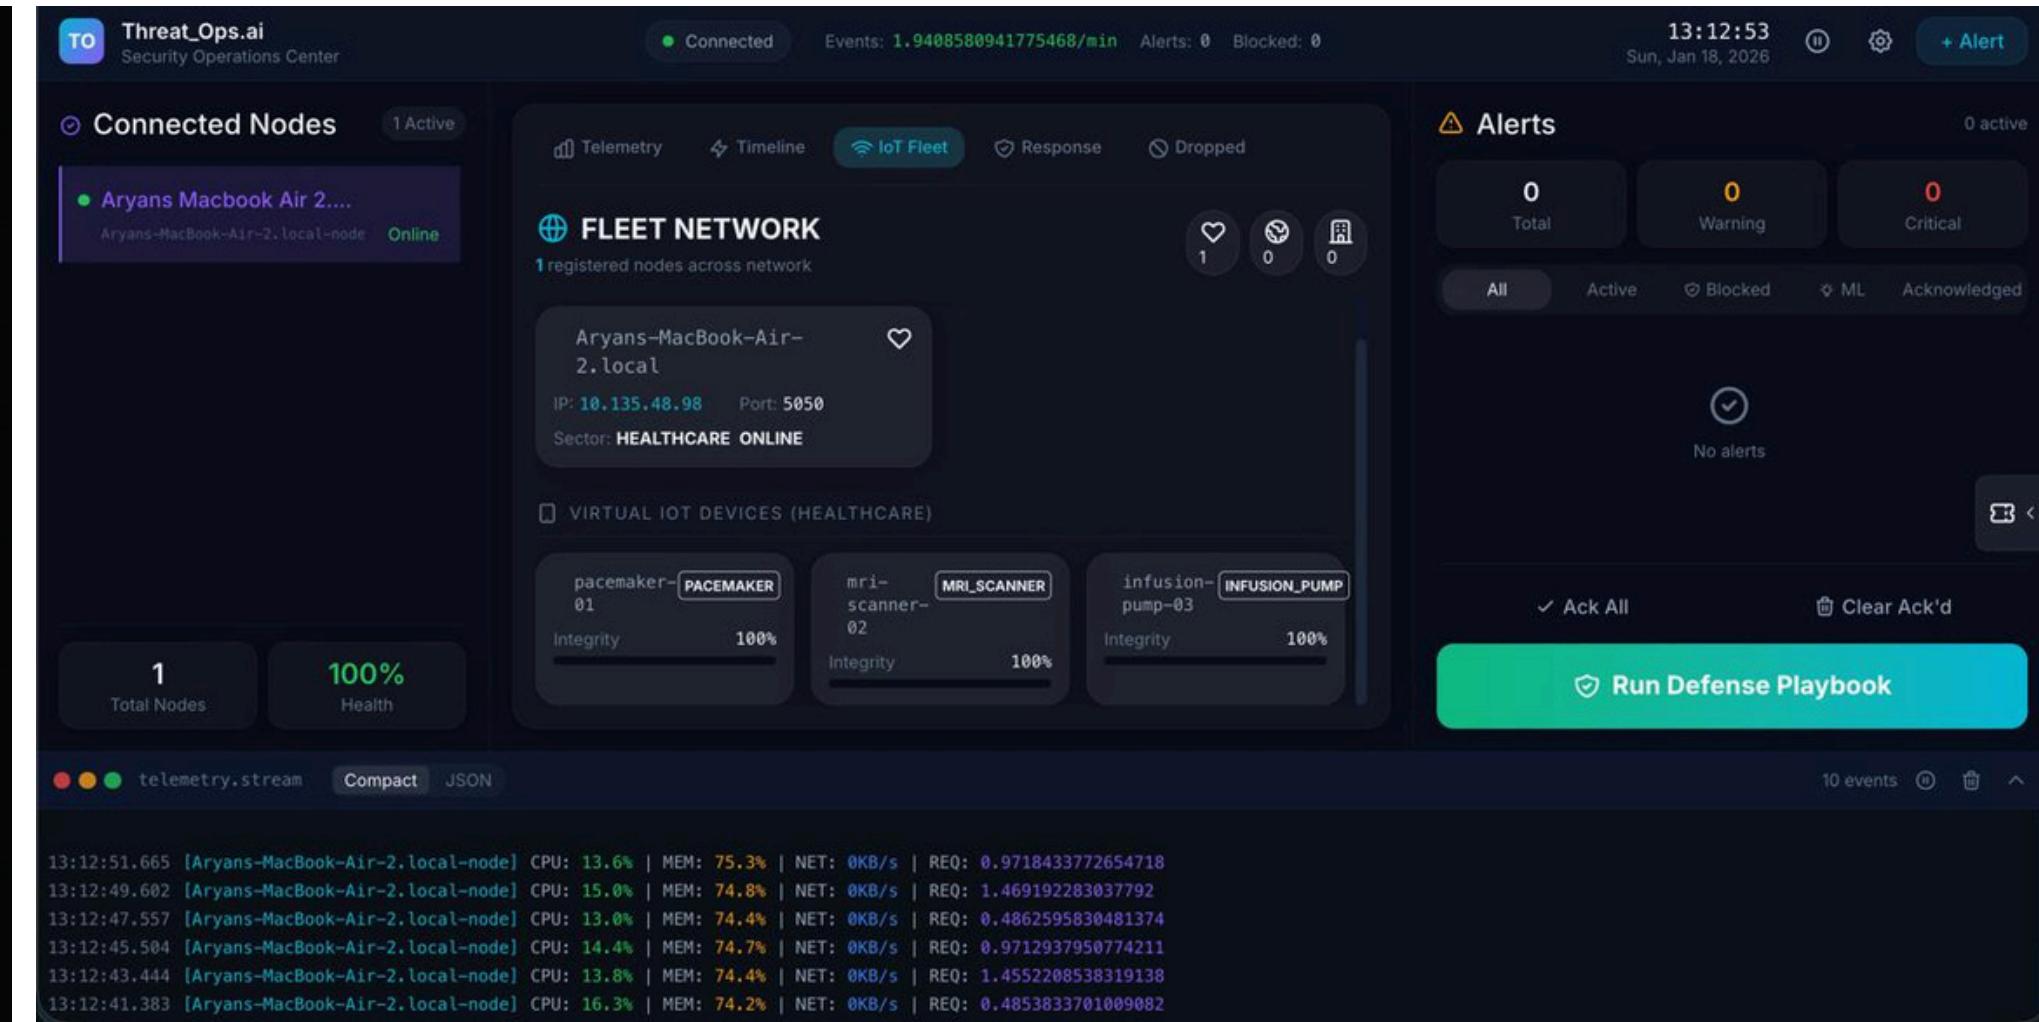
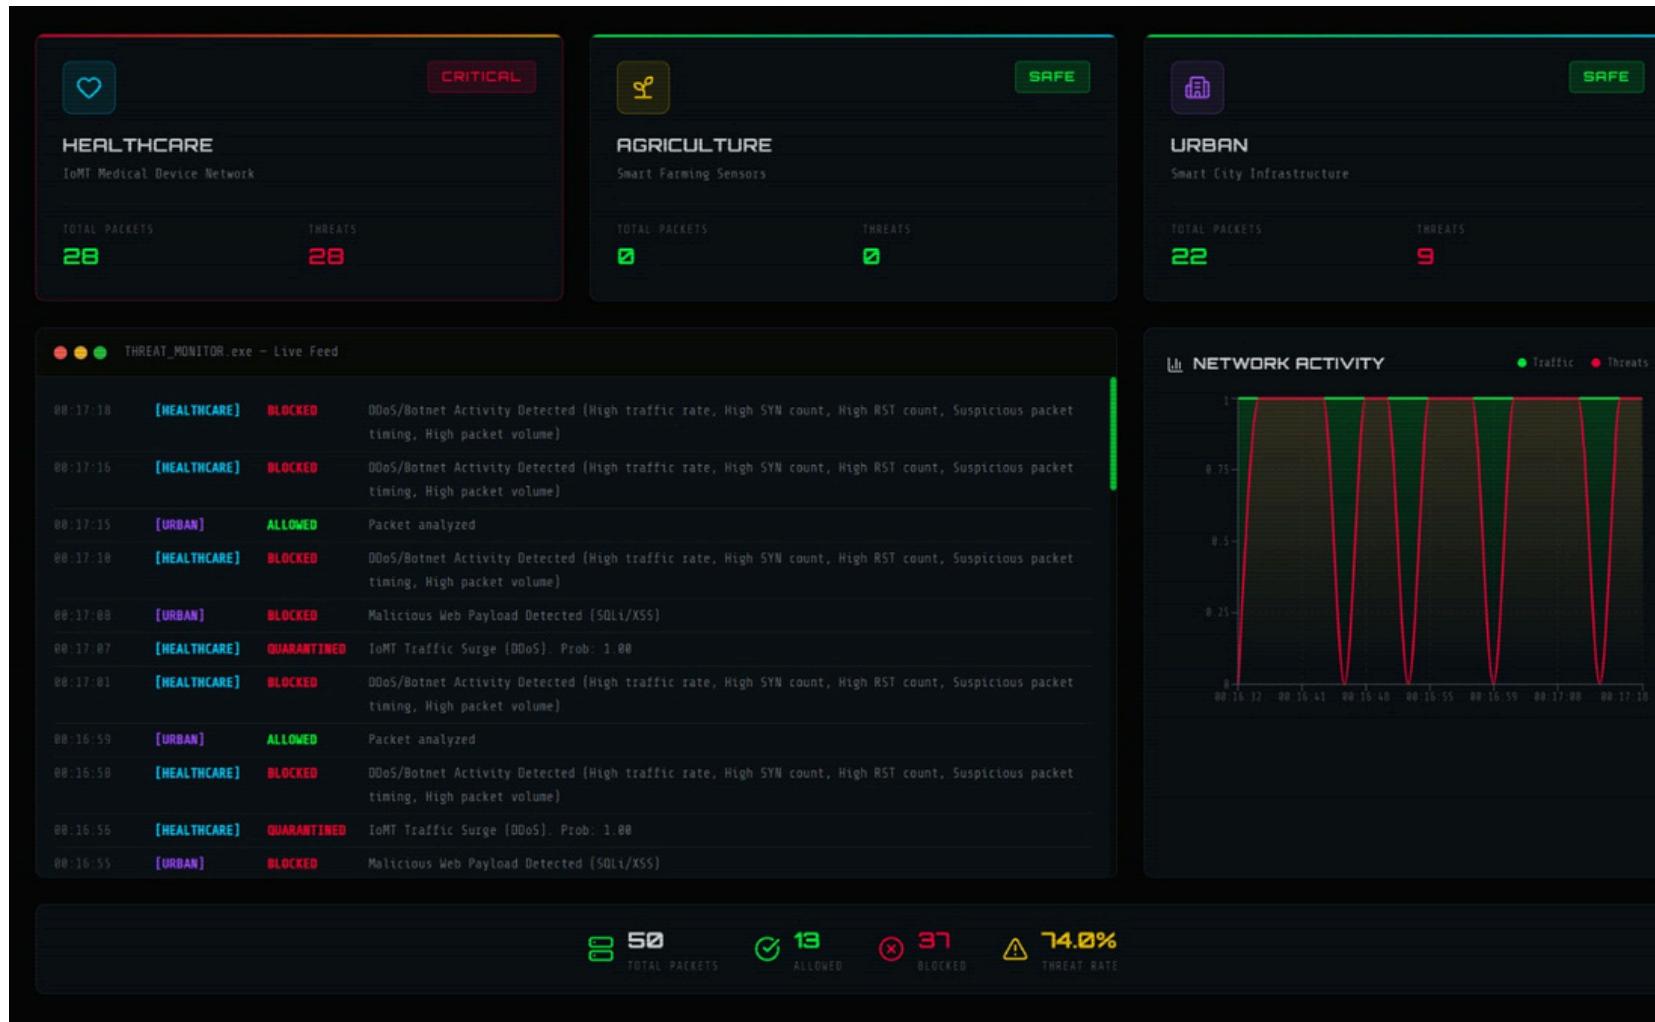
Logs: A log stream from `telemetry.stream` shows the following recent entries:

```
01:00:29.314 [Het-Workspace-node] CPU: 3.6% | MEM: 84.0% | NET: 250KB/s | REQ: 0
01:00:29.144 [Jays-MacBook-Air.local-node] CPU: 12.3% | MEM: 66.8% | NET: 0KB/s | REQ: 0
01:00:27.156 [Het-Workspace-node] CPU: 2.8% | MEM: 84.1% | NET: 252KB/s | REQ: 0
01:00:26.973 [Jays-MacBook-Air.local-node] CPU: 11.5% | MEM: 66.5% | NET: 0KB/s | REQ: 0
01:00:25.075 [Het-Workspace-node] CPU: 3.0% | MEM: 84.1% | NET: 252KB/s | REQ: 0
01:00:24.887 [Jays-MacBook-Air.local-node] CPU: 7.2% | MEM: 66.7% | NET: 0KB/s | REQ: 0
01:00:22.956 [Het-Workspace-node] CPU: 4.2% | MEM: 84.1% | NET: 253KB/s | REQ: 0
```

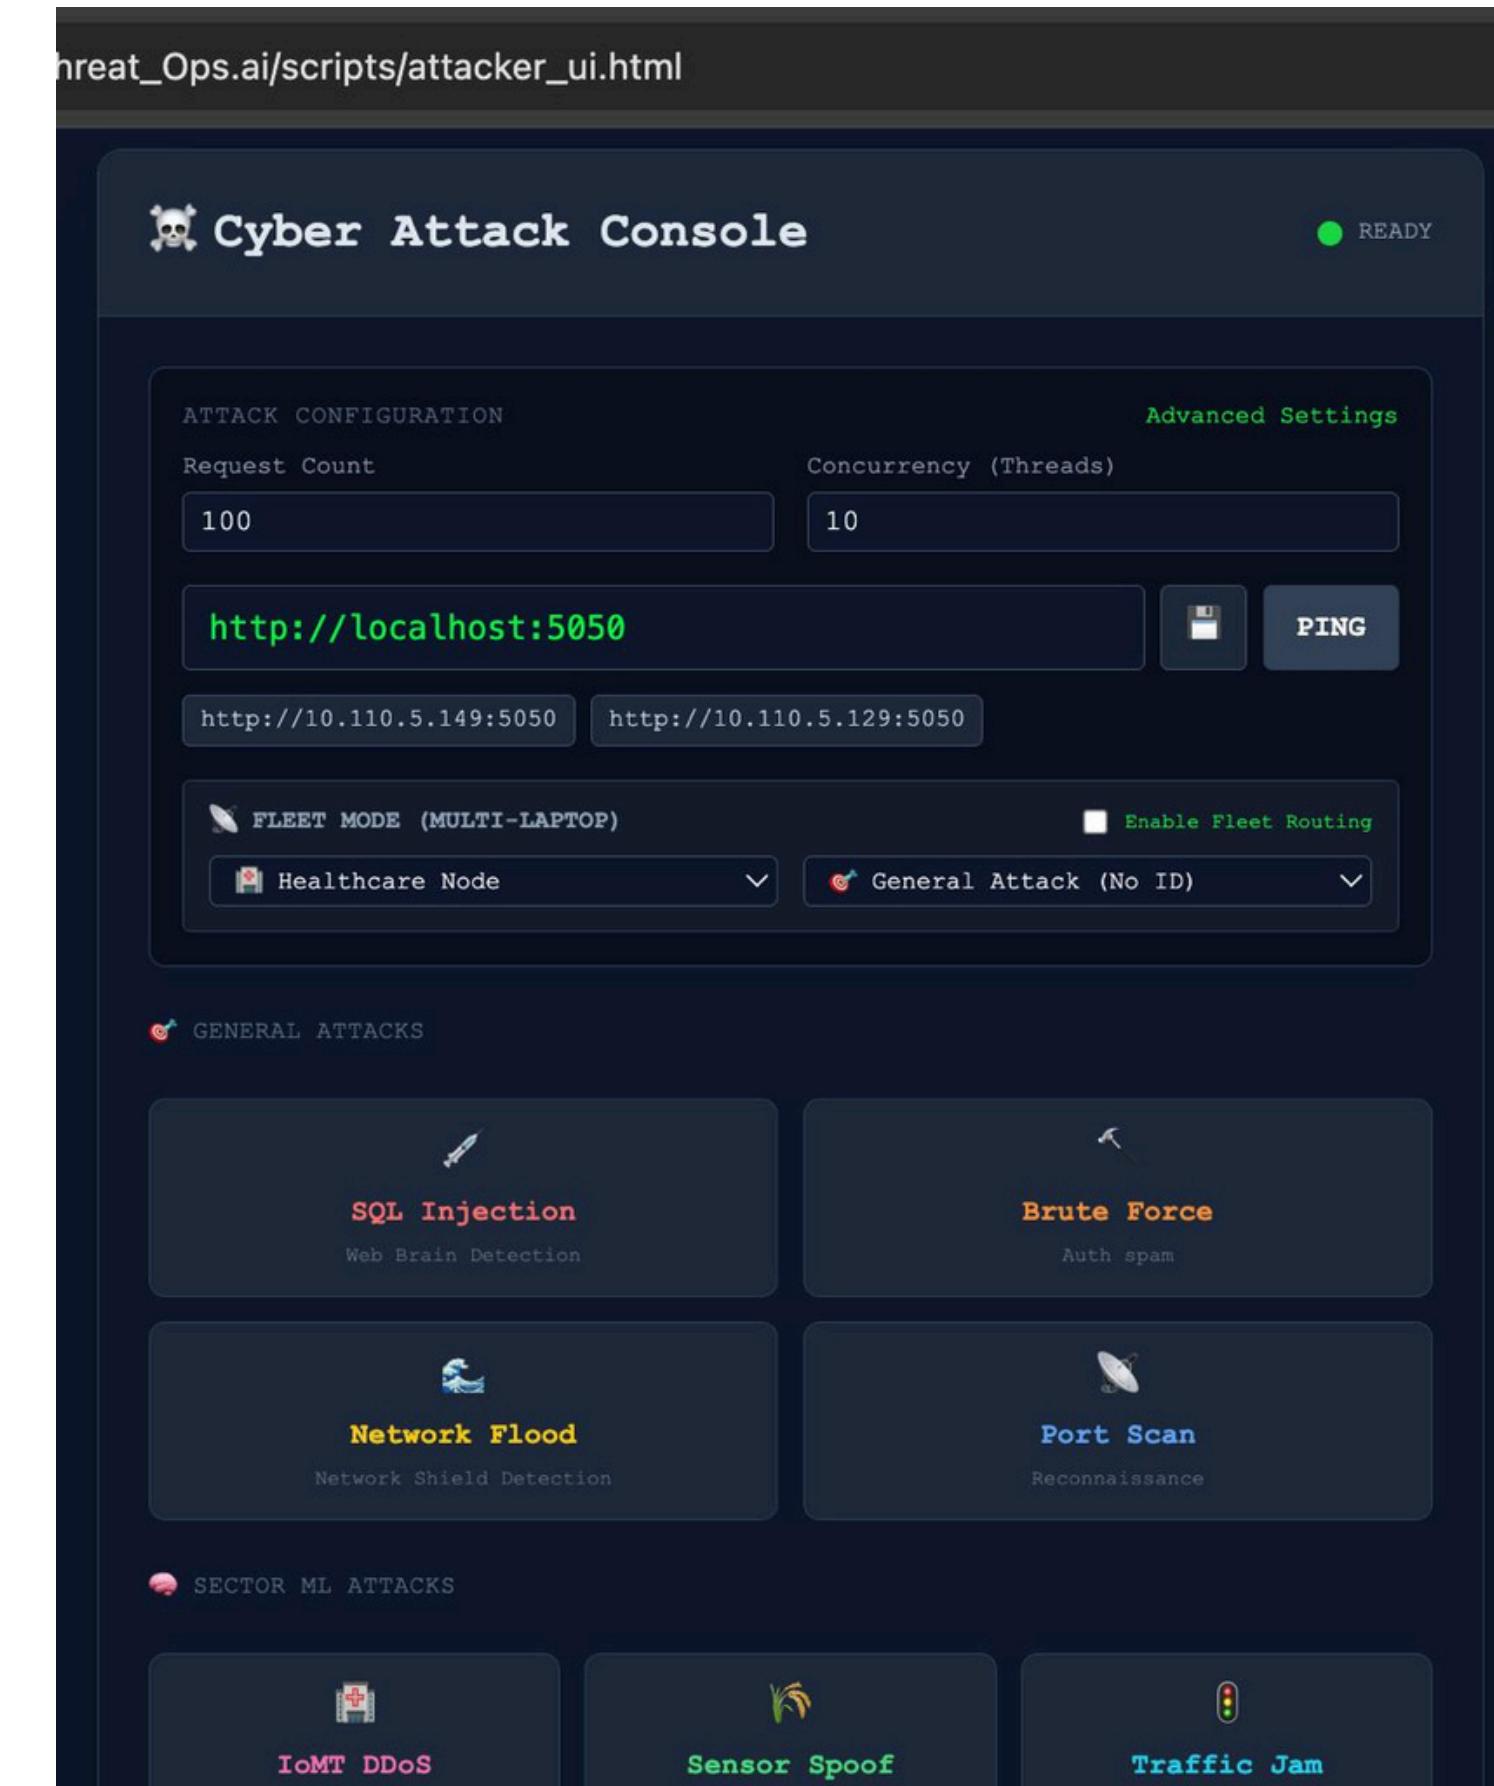
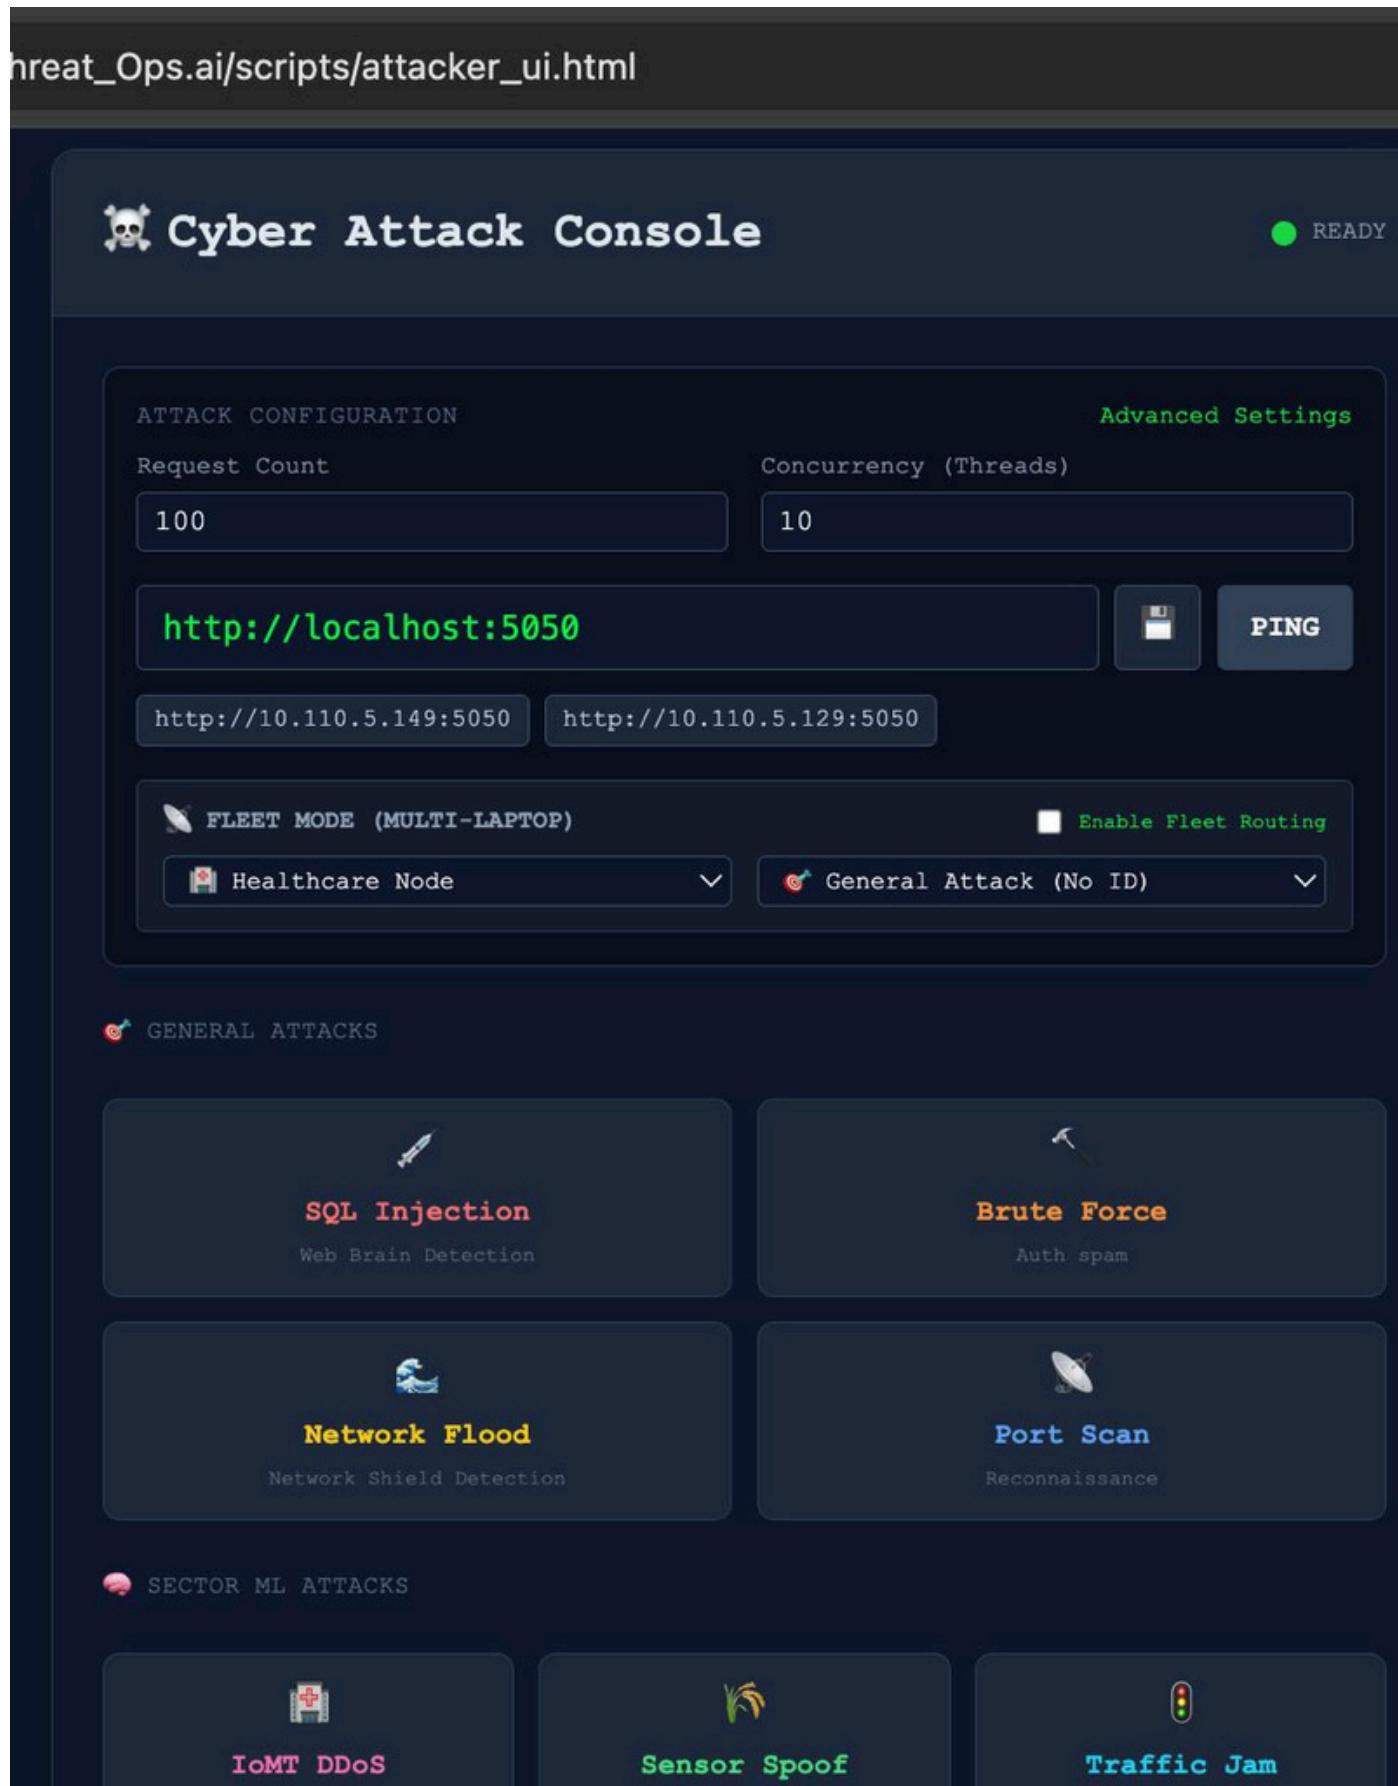
```
PS C:\Users\Het Ashishbhai Modi\idk> python .\monitor_server.py
 Threat_Ops Universal Agent
Device: Het-Workspace-node
IP: 10.110.5.149
Server: http://10.110.5.98:8001/ingest

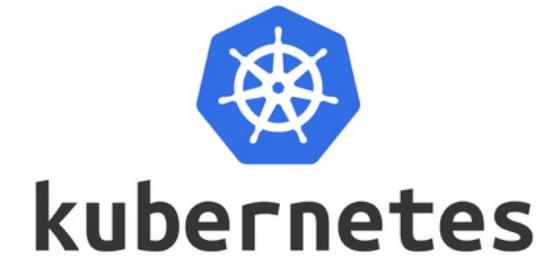
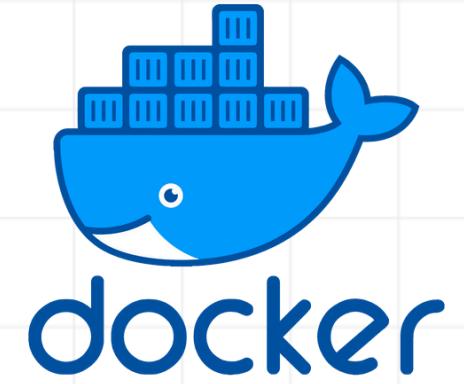
  Telemetry Agent started. Reporting to http://10.110.5.98:8001/ingest
* Serving Flask app 'monitor_server'
* Debug mode: off
WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
* Running on all addresses (0.0.0.0)
* Running on http://127.0.0.1:5050
* Running on http://10.110.5.149:5050
Press CTRL+C to quit
10.110.5.98 - - [18/Jan/2026 01:14:31] "OPTIONS /data HTTP/1.1" 200 -
⚠ [SECURITY] Query received: SELECT * FROM users WHERE id = 1 OR 1=1
⚡ Forwarded alert to Main Server: 200
10.110.5.98 - - [18/Jan/2026 01:14:31] "POST /data HTTP/1.1" 200 -
10.110.5.98 - - [18/Jan/2026 01:14:36] "OPTIONS /data HTTP/1.1" 200 -
⚠ [SECURITY] Query received: SELECT * FROM users WHERE id = 1 OR 1=1
⚡ Forwarded alert to Main Server: 200
10.110.5.98 - - [18/Jan/2026 01:14:37] "POST /data HTTP/1.1" 200 -
10.110.5.98 - - [18/Jan/2026 01:14:38] "OPTIONS /login HTTP/1.1" 200 -
⚠ [SECURITY] Login attempt for user: admin from 10.110.5.98
10.110.5.98 - - [18/Jan/2026 01:14:39] "OPTIONS /login HTTP/1.1" 200 -
⚠ [SECURITY] Login attempt for user: admin from 10.110.5.98
⚡ Forwarded alert to Main Server: 200
10.110.5.98 - - [18/Jan/2026 01:14:39] "POST /login HTTP/1.1" 401 -
⚡ Forwarded alert to Main Server: 200
10.110.5.98 - - [18/Jan/2026 01:14:39] "POST /login HTTP/1.1" 401 -
10.110.5.98 - - [18/Jan/2026 01:14:39] "OPTIONS /login HTTP/1.1" 200 -
⚠ [SECURITY] Login attempt for user: admin from 10.110.5.98
⚡ Forwarded alert to Main Server: 200
10.110.5.98 - - [18/Jan/2026 01:14:39] "POST /login HTTP/1.1" 401 -
10.110.5.98 - - [18/Jan/2026 01:14:39] "OPTIONS /login HTTP/1.1" 200 -
⚠ [SECURITY] Login attempt for user: admin from 10.110.5.98
⚡ Forwarded alert to Main Server: 200
10.110.5.98 - - [18/Jan/2026 01:14:39] "POST /login HTTP/1.1" 401 -
10.110.5.98 - - [18/Jan/2026 01:14:40] "OPTIONS /login HTTP/1.1" 200 -
```

Backend server

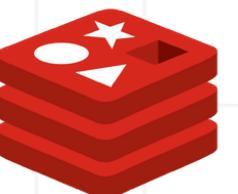
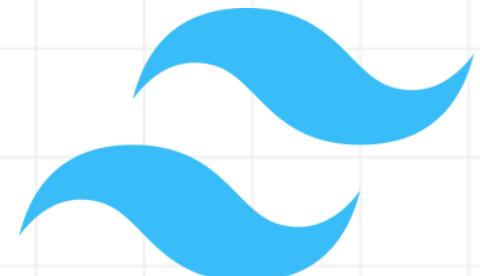


ATTACK SIMULATION

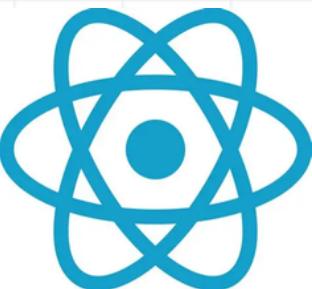




TECHNOLOGIES USED



redis



P Y T O R C H



RESOURCES

DATASET LINKS :

- [HTTPS://WWW.KAGGLE.COM/DATASETS/ETS/PETERFRIEDRICH1/CICDARKNET2020-INTERNET-TRAFFIC](https://www.kaggle.com/datasets/peterfriedrich/cicdarknet2020-internet-traffic)
- [HTTPS://DATA.MENDELEY.COM/DATASETS/2BW34GHT8B/2](https://data.mendeley.com/datasets/2bw34ght8b/2)
- [HTTPS://WWW.KAGGLE.COM/DATASETS/ETS/WISAM1985/IOT-AGRICULTURE-2024](https://www.kaggle.com/datasets/wisam1985/iot-agriculture-2024)
- [HTTPS://WWW.KAGGLE.COM/DATASETS/ETS/HIMADRI07/CICIOT2023/DATA](https://www.kaggle.com/datasets/himadri07/ciciot2023/data)
- [HTTPS://WWW.KAGGLE.COM/DATASETS/ETS/EVG3N1J/HTTPPARAMSDATABASE](https://www.kaggle.com/datasets/evg3n1j/httpparamsdatabase)